


Agentic AI Governance Drift Cone™

INDEPENDENT GOVERNANCE & ASSURANCE METHODOLOGY

A structured approach to identify, assess and assure the governance state of agentic AI systems and ecosystems.
From intent to assurance — and beyond.

Adnan Shamim

Founder of
The First Agentic AI
Audit Ecosystem™



AGENTIC AI GOVERNANCE DRIFT CONE FRAMEWORK™

V1.0

Independent Analytical, Governance, Control-Assessment & Assurance Methodology for Agentic AI

Version: 1.0

Status: Foundational Methodology Edition

Purpose: Governance Drift Analysis, Control Assessment & Assurance

Scope: Agentic AI Systems and Ecosystems

1. PURPOSE

The **Agentic AI Governance Drift Cone Framework™** provides a structured methodology for identifying, analyzing, assessing, testing, evidencing, reporting, and continuously reassessing differences between intended governance and the actual state of an Agentic AI system.

The Framework examines governance across the progression:

DOCUMENTED

→ CRITERIA-BOUND

→ AUTHORIZED

→ BOUNDED

→ ENFORCED WHERE REQUIRED

→ IMPLEMENTED

→ OPERATING

→ EVIDENCED

→ TESTED

→ EFFECTIVE

→ DEMONSTRABLE

→ REVALIDATED WHEN REQUIRED



The Framework is designed to support governance, risk, internal audit, compliance, control assessment, independent assurance, and related functions.

It does not itself establish legal or regulatory requirements.


2. STATUS OF THE FRAMEWORK

This Framework is an analytical and methodological framework.

It is not, by itself:

- a law;
- regulation;
- regulatory direction;
- regulatory interpretation;

- legal opinion;
- professional standard;
- accounting standard;
- auditing standard;
- certification scheme;
- licensing requirement;
- contractual obligation; or




- statutory requirement.

Use of this Framework does not, by itself, establish compliance with any external requirement.

3. REGULATORY, LEGAL & PROFESSIONAL NON-SUPREMACY PRINCIPLE

No provision of this Framework shall be interpreted as overriding, replacing, weakening, waiving, modifying, or superseding any applicable:

- law or statute;
- regulation;
- regulatory or supervisory requirement;
- governmental or judicial order;



WHERE AN APPLICABLE AUTHORITATIVE REQUIREMENT DIFFERS FROM THIS FRAMEWORK, THE APPLICABLE AUTHORITATIVE REQUIREMENT PREVAILS.

4. NO LEGAL OR REGULATORY DETERMINATION

Application of this Framework shall not, without separate appropriate authority and criteria, be represented as:

- legal advice;

- legal interpretation;
- regulatory interpretation;
- regulatory approval;
- regulatory certification;
- statutory compliance certification;
- professional certification;
- legal compliance determination; or
- evidence that a regulator, court, professional body, or other competent authority accepts the Framework.

Where legal or regulatory interpretation is required, the matter shall be referred to the appropriately authorized legal, regulatory, compliance, or professional authority.

5. JURISDICTIONAL NEUTRALITY

The Framework is jurisdiction-neutral unless specific jurisdictional requirements are expressly adopted as applicable assessment criteria.

References to laws, regulations, standards, frameworks, supervisory expectations, or professional requirements shall not automatically make those requirements applicable.

APPLICABILITY MUST BE ESTABLISHED SEPARATELY.

6. CRITERIA-FIRST ASSURANCE RULE

No conclusion shall be formed merely because an Agentic AI system conforms to this Framework.

The applicable assessment shall establish, as appropriate:

ASSERTION

→ CRITERIA

→ AUTHORITY

→ SCOPE

→ PERIOD

→ RISK

→ EVIDENCE


- TESTING
- LIMITATIONS
- CONCLUSION

NO CONCLUSION SHALL EXCEED THE CRITERIA AGAINST WHICH IT WAS ASSESSED.

7. FUNDAMENTAL AGENTIC AI GOVERNANCE PRINCIPLE

AI CAPABILITY DOES NOT CREATE AUTHORITY.

Therefore:



Each state requires its own appropriate basis.

9. GOVERNANCE DRIFT

Definition

Governance Drift is a material or potentially material difference between an applicable intended, authorized, required, implemented, observed, or historical governance state.

Drift is an analytical finding, not automatically a control failure.

Therefore:

DETECTED DIFFERENCE ≠ CONTROL FAILURE ≠ ASSURANCE QUALIFICATION

A difference shall be assessed against applicable criteria, risk, materiality, control objectives, evidence, and circumstances before being classified.

Agentic AI Governance Drift Cone™
INDEPENDENT GOVERNANCE & ASSURANCE METHODOLOGY
A structured approach to defining, assessing and assuring the governance of agentic AI systems and ecosystems.
From intent to assurance — and beyond.

Adnan Shamim
The First Agentic AI
Audit Ecosystem™

10. FIVE-STATE GOVERNANCE MODEL

The Framework distinguishes:

EXPECTED STATE

What governance is intended or required to achieve.

AUTHORIZED STATE

What the applicable authority has actually permitted within defined conditions.

IMPLEMENTED STATE

What has been technically or procedurally established.

OBSERVED STATE

What evidence indicates actually occurred or operated.

HISTORICAL STATE

What the evidence establishes occurred during a prior relevant period.

Therefore:

EXPECTED ≠ AUTHORIZED ≠ IMPLEMENTED ≠ OBSERVED ≠ HISTORICAL

Conflating these states can create false assurance.

11. AUTHORITY BOUNDARY STANDARD

Every consequential Agentic AI authority should be capable of being evaluated across:

WHO — authorized actor or accountable principal

WHAT — permitted action

WHY — authorized purpose

WHERE — permitted environment/system

WHEN — validity period

HOW — permitted mechanism

HOW MUCH — quantitative limit

DATA — permitted information boundary

TOOLS — permitted capabilities

CONDITIONS — preconditions

APPROVAL — required human or organizational approval

ESCALATION — mandatory escalation conditions

PROHIBITIONS — expressly prohibited actions

DELEGATION — permitted delegation conditions

EXPIRY — authority expiration

REVOCACTION — revocation conditions and mechanism


12. AUTHORITY VALIDITY

Authority is valid only if it is bounded, enforceable, and still governs when it acted?

Authorization establishes a permitted condition only within its valid boundary and applicable conditions.

Authority may become invalid, restricted, suspended, expired, revoked, or otherwise inapplicable because of:


- time;
- context;
- policy;
- risk;
- identity;
- permissions;
- scope;



- data conditions;
- system state;
- approval status;
- material change;
- emergency conditions; or
- revocation.

Therefore:

AUTHORIZATION AT DECISION TIME DOES NOT AUTOMATICALLY PROVE AUTHORIZATION AT EXECUTION TIME.



Failure of a mandatory condition shall prevent execution where preventive enforcement is required.

NO VALID AUTHORIZATION → NO AUTHORIZED EXECUTION

14. ENFORCEABILITY

A documented boundary that cannot be appropriately enforced shall not be represented as an effective preventive control merely because the boundary exists in policy.

Where alternative controls exist, their effectiveness must be assessed against their own defined objectives and criteria.

If an unauthorized path remains technically reachable and materially capable of producing a prohibited outcome, the existence of an authorized path does not by itself demonstrate preventive enforcement.

15. GOVERNANCE DRIFT DOMAINS

Governance Drift may manifest through:

1. AUTHORITY DRIFT

Authority no longer corresponds to approved authorization.

A structured approach to identify, assess and assure the governance state of agentic AI systems and ecosystems, beyond.

2. EVIDENCE DRIFT

Evidence no longer reliably represents the relevant governance or execution state.

3. MODEL / BEHAVIOURAL DRIFT

System behaviour changes materially relative to applicable governance assumptions.

4. ACCOUNTABILITY DRIFT


Responsibility becomes unclear, displaced, or inconsistent with actual decision rights.

5. BOUNDARY DRIFT

System capability or operating scope expands beyond the approved boundary.

6. RESILIENCE DRIFT

Controls degrade during failure, overload, outage, dependency failure, or adverse conditions.



16. DRIFT ASSESSMENT

For every material detected difference, assess:

1. What changed?
2. Which state changed?
3. What criteria apply?
4. Is the difference authorized?

5. Is it material?
6. What risk does it create?
7. Which control objective is affected?
8. What evidence supports the conclusion?
9. Was enforcement maintained?
10. Was accountability maintained?
11. Was revocation maintained?
12. Does the difference constitute a control deficiency?
13. Does it affect the assurance conclusion?

Agentic AI Governance Drift Cone™

A structured approach to identify, assess and assure the governance state of agentic AI systems and ecosystems.
From intent to assurance — and beyond.

Adnan Shamim

Founder of
The First Agentic AI
Audit Ecosystem™

17. CHANGE GOVERNANCE

Material changes to an Agentic AI system should be assessed for impact on:

- authority;

- permissions;

- policies;

- authorized,

- bounded, enforceable,

- and governed

- when it acted?

- tools;

- integrations;

- data;

- workflows;


- thresholds;

- approval requirements;

- human oversight;

- monitoring;

- evidence;



- revocation;
- resilience; and
- assurance conclusions.

Prior assurance shall not automatically be assumed to remain valid following a material change.

18. EMERGENCY CHANGE CONTROL

Emergency changes may require accelerated procedures.

However:

EMERGENCY DOES NOT AUTOMATICALLY MEAN EXEMPT FROM CONTROL.

A structured approach to identify, assess and assure the governance state of agentic AI systems and ecosystems.

Where normal controls are bypassed or modified, the exception, authorization, duration, compensating controls, evidence, accountability, and post-event review should be appropriately documented.

19. MODEL & BEHAVIOURAL DRIFT

Governance assessment shall consider material changes in:

- model behaviour;
- model version;

- system prompts;

- retrieval sources;

- tools;

- agent policies;

- decision patterns;



- confidence;

- error rates;

- hallucination behaviour;

- data distribution;

- environment;




- downstream dependencies.

A model remaining technically available does not establish that its previously established assurance remains valid.

20. DATA AUTHORITY BOUNDARY

Data access shall be evaluated independently from action authority.

ACCESS TO DATA DOES NOT CREATE AUTHORITY TO USE, DISCLOSE, MODIFY, TRANSFER, OR ACT UPON THAT DATA.



21. THIRD-PARTY & DEPENDENCY GOVERNANCE

Third-party models, APIs, platforms, tools, datasets, infrastructure, plugins, and external services may form part of the effective control environment.

Where relevant, assessment should consider:

- dependency;
- availability;
- security;

- authorization;
- contractual obligations;
- data boundaries;
- change management;
- monitoring;
- evidence;
- failure modes;
- termination/revocation.

Agentic AI Governance Drift Cone™

Adnan Shamim

Founder of
AI Governance
Audit Ecosystem™

Where evidence is unavailable because it is outside the assurance boundary, that limitation shall be explicitly recognized.



CORE ASSURANCE TEST

- relevance;
- sufficiency;
- appropriateness;
- reliability;

NON-NEGOTIABLE PRINCIPLES

- 1 Capability ≠ Authority ≠ Execution
- 2 Documentation ≠ Effectiveness
- 3 Authorization ≠ Execution Authorization
- 4 Evidence Validity ≠ Evidence Completeness
- 5 No Evidence ≠ Evidence of No Event
- 6 Automation ≠ Transfer of Accountability
- 7 Material Change → Revalidation

GOVERNANCE DRIFT

Governance states are distinct — not interchangeable.

E > A > I > O > H

EXPECTED AUTHORIZED IMPLEMENTED OBSERVED HISTORICAL

- authenticity;
- provenance;

PROVE THE GOVERNANCE STATE.
PROVE THE AUTHORITY.
PROVE THE CONTROL.
PROVE THE EVIDENCE.
DO NOT OVERSTATE ASSURANCE.



Applicable law, regulation, contractual and professional requirements prevail.

- completeness;
- integrity;
- chronology;
- accessibility;
- independence where required.

A technically valid record may still be insufficient for the assertion being tested.

23. EVIDENCE COMPLETENESS

Three separate questions shall be maintained:

UNIVERSE COMPLETENESS

Did the assessment identify the relevant population?

STATE COMPLETENESS

Did the assessment capture relevant state changes?

EVIDENCE COMPLETENESS

Does the available evidence sufficiently cover the relevant assertion and period?

Therefore:

IDENT GOVERNANCE & ASSURANCE METHODOLOGY

A structured approach to identify, assess and assure the governance state of agentic AI systems and ecosystems.

EVIDENCE VALIDITY ≠ EVIDENCE COMPLETENESS

24. NEGATIVE-ASSURANCE RULE

The following shall not be treated as equivalent:

No evidence detected

≠ evidence that no event occurred

No detected change

≠ proof that no change occurred

No logged event

≠ proof that no event occurred

No observed failure

≠ proof of effectiveness


No captured act

≠ proof that no act occurred

Insufficient evidence shall not be converted into a positive assurance conclusion.

25. AGENT EXECUTION EVIDENCE

For consequential actions, evidence should permit reconstruction, as applicable, of:



WHO

→ WHAT

→ WHY

→ AUTHORITY

→ POLICY

→ CONTEXT

→ APPROVAL

→ DECISION

→ EXECUTION

→ RESULT

→ EXCEPTION

→ REVOCATION / TERMINATION

Agentic AI Governance Drift Cone™

INDEPENDENT GOVERNANCE & ASSURANCE METHODOLOGY

The required evidence depends on the applicable assertion, risk, criteria, scope, and engagement objective.

Adnan Shamim

Founder of
The First Agentic AI

26. HUMAN OVERSIGHT

Human involvement shall not automatically be treated as effective human control.

Distinguish:

VISIBILITY ≠ REVIEW ≠ APPROVAL ≠ CONTROL

A human approval mechanism shall not be treated as an effective control where the relevant agent or an alternative execution path can materially bypass the approval requirement.

27. DEMONSTRABLE REVOCABILITY

Revocation shall be assessed as a control capability, not merely as a documented instruction.

A robust revocation mechanism should, where applicable, be:

- authorized;
- enforceable;
- timely;
- fail-closed where required;
- independently testable;

- traceable;
- integrity-protected;
- resistant to bypass;
- capable of preventing further prohibited execution.

A STOP COMMAND IS AN INSTRUCTION.

DEMONSTRATED INABILITY TO PERFORM THE PROHIBITED ACTION IS EVIDENCE RELEVANT TO REVOCATION EFFECTIVENESS.

28. CONTROL-UNIVERSE COMPLETENESS

INDEPENDENT GOVERNANCE & ASSURANCE METHODOLOGY

A structured approach to identify, assess and verify the controls that govern AI systems and ecosystems, from intent to assurance and beyond.

Adnan Shamim
 Founder of
 The First Agentic AI
 Audit Ecosystem™

The assessment should consider:

- primary controls;
- compensating controls;

CORE ASSURANCE TEST

- preventive controls;
- detective controls;
- corrective controls;
- manual controls;
- automated controls;
- inherited controls;

NON-NEGOTIABLE PRINCIPLES

1 Capability ≠ Authority ≠ Execution 2 Documentation ≠ Effectiveness 3 Authorization ≠ Execution Authorization 4 Evidence Validity ≠ Evidence Completeness	5 No Evidence ≠ Evidence of No Event 6 Automation ≠ Transfer of Accountability 7 Material Change → Revalidation
---	---

GOVERNANCE DRIFT

Governance states are distinct — not interchangeable.

E

>

A

>

I

>

O

>

H

EXPECTED AUTHORIZED IMPLEMENTED OBSERVED HISTORICAL

GOVERNANCE STATE.
 PROVE THE AUTHORITY.
 PROVE THE CONTROL.
 PROVIDE THE EVIDENCE.
 DO NOT OVERSTATE ASSURANCE.

Applicable law, regulation, contractual and professional requirements prevail.

- third-party controls;
- dependencies;
- exception controls;
- fallback mechanisms;
- alternative execution paths.

29. RISK & MATERIALITY

Not every governance difference has equal significance.

Assessment should consider:

LIKELIHOOD

- × IMPACT
- × EXPOSURE
- × CONTROL DEPENDENCY
- × REVERSIBILITY
- × DETECTABILITY
- × DURATION

Materiality shall be determined against the applicable engagement criteria and organizational context rather than by this Framework alone.


30. EXCEPTION GOVERNANCE

An exception should identify, as applicable:

- authority;
- rationale;
- scope;
- Was the AI authorized, bounded, enforceable and still governed where it started?
- affected systems;
- affected actions;

- risk;
- compensating controls;

- start date;
- expiry;
- accountable owner;
- approval;
- monitoring;
- remediation;
- closure evidence.



An exception shall not silently become permanent operating authority.

31. LIMITATION GOVERNANCE

A limitation shall be explicitly identified where relevant evidence, systems, populations, periods, dependencies, or testing are unavailable or restricted.

Limitations shall not be hidden within an overall conclusion.

A LIMITATION SHALL NOT BE CONVERTED INTO AN EFFECTIVENESS CLAIM.

32. CONTROL FAILURE DECISION LOGIC

A detected difference should be evaluated through:

DIFFERENCE DETECTED

→ **CRITERIA IDENTIFIED**

→ **CONTROL OBJECTIVE IDENTIFIED**

→ **RISK ASSESSED**

→ **MATERIALITY ASSESSED**

→ **CONTROL DESIGN ASSESSED**

→ **IMPLEMENTATION ASSESSED**

→ **OPERATING EFFECTIVENESS TESTED**

→ **EVIDENCE EVALUATED**

→ **CONCLUSION FORMED**

Only where the applicable basis supports it should a difference be classified as a control deficiency or failure.

33. ACCOUNTABILITY

Accountability shall remain with the person, function, committee, management body, system owner, or other entity having the applicable mandate.

Agent autonomy does not automatically transfer organizational accountability to the AI system.

AUTOMATION DOES NOT EXTINGUISH ACCOUNTABILITY.

34. SEGREGATION OF DUTIES

Where required by applicable criteria, the Framework should evaluate whether incompatible responsibilities are appropriately separated, including:

- authorization;
- configuration;
- execution;
- monitoring;
- exception approval;
- evidence administration;
- assurance.

The appropriate segregation model depends on the applicable control objective and organizational structure.

35. INDEPENDENT ASSURANCE

Where the Framework is used in an independent assurance engagement, its application shall preserve applicable requirements relating to:

- independence;
- bounded, enforceable, and objective when it acted?
- competence;

- due professional care;
- confidentiality;

- evidence;
- documentation;
- professional judgment.

The Framework shall not be used to predetermine the assurance conclusion.

36. ASSURANCE ARCHITECTURE

The assessment should distinguish:

- evidence;
- documentation;
- professional judgment.

GOVERNANCE INTENT

- AUTHORITY
- BOUNDARY
- ENFORCEMENT
- IMPLEMENTATION
- OPERATION
- EVIDENCE
- TESTING
- REVOCATION
- ASSURANCE
- CONTINUOUS REVALIDATION

Agentic AI Governance Drift Cone™
Failure at one layer shall not automatically be inferred from failure at another.

A structured approach to identify, assess and assure the governance state of agentic AI systems and ecosystems.
Prove intent, demonstrate, and improve.

Adnan Shamim

Founder of
The First Agentic AI
Audit Ecosystem™

37. ASSURANCE CONCLUSION HIERARCHY

Subject to applicable professional requirements and engagement criteria, conclusions may include:

DEMONSTRABLY EFFECTIVE

Evidence and testing support the applicable assertion within the defined boundary.

EFFECTIVE WITH LIMITATIONS

The assertion is supported, subject to identified limitations that do not invalidate the conclusion.

PARTIALLY EFFECTIVE

Material aspects of the applicable control objective are supported, while deficiencies remain.

INEFFECTIVE

Evidence and testing demonstrate that the applicable control objective is not achieved.

UNABLE TO CONCLUDE


Evidence or access limitations prevent a sufficiently supported conclusion.

The terminology may be adapted where applicable professional requirements prescribe different conclusion structures.

38. ASSURANCE BOUNDARY PRINCIPLE

Every assurance conclusion shall define:

- subject matter;
- assertion;
- criteria;
- systems;
- populations;
- period;
- jurisdictions where relevant;
- dependencies;
- exclusions;



Examples include:


- material policy change;
- authority change;
- model change;
- system change;
- tool change;
- data change;
- integration change;
- risk change;

- regulatory change;
- material incident;
- control failure;
- revocation failure;
- material behavioural drift.

40. REAUTHORIZATION TRIGGERS

Reauthorization or reassessment should be considered when applicable conditions materially change.

Potential triggers include:



41. PROFESSIONAL & REGULATORY COMPATIBILITY

This Framework is intended to operate **alongside**, not instead of, applicable professional and regulatory requirements.

Where an engagement is subject to mandatory professional standards, those standards govern the engagement.

Where a regulator establishes specific requirements, those requirements govern the applicable regulatory assessment.

Where legal interpretation is required, qualified legal counsel or the relevant competent authority shall determine the applicable legal position.

This Framework may provide additional analytical structure only to the extent that such use does not conflict with the governing requirements.

42. NO AUTOMATIC COMPLIANCE EQUIVALENCE

Conformance with this Framework shall not be described as equivalent to:

- legal compliance;
- regulatory compliance;
- regulatory approval;
- regulatory certification;
- professional certification;
- audit certification;
- conformity with an external standard; or
- immunity from legal, regulatory, operational, financial, security, privacy, or professional



Any such claim requires separate, criteria-specific substantiation.

43. NO SAFE-HARBOUR CLAIM

Implementation of this Framework does not create a safe harbour, exemption, immunity, defence, regulatory acceptance, or legal protection.

The existence of:

- a policy;
- a control;
- a test;
- an evidence record;
- an assurance report; or



- a Framework implementation

does not eliminate the possibility of:

- error;
- fraud;
- unauthorized activity;
- control failure;
- non-compliance;
- litigation;
- regulatory action;
- operational failure; or
- other adverse outcomes.



44. CONFLICT ESCALATION

Where an apparent conflict exists between this Framework and an applicable authoritative requirement:

1. identify the conflicting provisions;
2. identify the applicable authority;
3. determine the relevant jurisdiction;
4. document the conflict;

5. apply the authoritative requirement;
6. limit, modify, or exclude the conflicting Framework provision; and
7. escalate unresolved material uncertainty to the appropriately authorized authority.

NO FRAMEWORK PROVISION SHALL BE USED TO JUSTIFY NON-COMPLIANCE WITH AN APPLICABLE AUTHORITATIVE REQUIREMENT.

45. INTERPRETATION SAFEGUARD

Where a provision could reasonably have multiple interpretations, the interpretation shall favor:

- preservation of applicable law;
- preservation of binding requirements;
- preservation of professional obligations;
- preservation of independence;
- prevention of unauthorized AI action;
- avoidance of assurance overstatement;
- transparency of limitations;
- clear accountability; and
- adherence to the defined assurance boundary.

Agentic AI Governance Drift Cone™

A structured approach to identify, assess and assure the governance state of agentic AI systems and ecosystems.
From intent to assurance — and beyond.

Adnan Shamim

Founder of
The First Agentic AI
Audit Ecosystem™

46. SEVERABILITY

If any provision of this Framework is determined to be invalid, unenforceable, impermissible, or inconsistent with an applicable authoritative requirement, that provision shall be limited or disregarded to the extent necessary, without invalidating the remaining provisions, to the extent permitted by applicable law.

47. NON-NEGOTIABLE GOVERNANCE RULES

1. AI capability does not create authority.

2. Capability ≠ recommendation ≠ approval ≠ authority ≠ execution.

3. Expected ≠ authorized ≠ implemented ≠ observed ≠ historical.

4. Detected difference ≠ control failure.

5. Control failure ≠ automatic assurance qualification without criteria.

6. Authorization must have a defined boundary.

7. Authority must remain valid when relevant.

8. Authority must be enforceable where preventive control is required.

9. Unauthorized paths must be considered.

10. Approval must not be confused with enforcement.

11. Human visibility is not automatically human control.

12. Evidence validity is not evidence completeness.

13. No evidence is not evidence of no event.

14. No log is not proof that no event occurred.

15. No observed failure is not proof of effectiveness.

16. Revocation must be demonstrable where required.

17. Exceptions require defined authority and duration.

18. Limitations must be disclosed.

19. Assurance must be criteria-bound.

20. Assurance must be scope-bound.

21. Assurance must be evidence-bound.

22. Assurance must be period-bound.

23. Assurance must not exceed its evidence.

24. Material change may invalidate prior assurance.

25. Third-party dependencies must be considered where relevant.

26. Model behaviour must be monitored where material.

27. Data authority must be assessed independently.

28. Accountability cannot be delegated merely because execution is automated.

29. Independence must be preserved.

30. Regulatory requirements prevail where applicable.

31. Legal requirements prevail where applicable.


32. Mandatory professional requirements prevail where applicable.

33. This Framework does not create legal authority.

34. This Framework does not create regulatory authority.

35. This Framework does not provide certification.

36. This Framework does not provide a safe harbour.



37. This Framework does not guarantee compliance.
38. This Framework does not guarantee effectiveness.
39. This Framework does not eliminate residual risk.
40. No claim shall exceed the authority, criteria, evidence, testing, scope, period, limitations, and professional basis supporting it.

48. ULTIMATE GOVERNANCE ASSURANCE TEST

An Agentic AI governance arrangement should not be represented as demonstrably effective unless the applicable assessment can establish, to the extent relevant to its defined scope:

1. **WHAT** was intended or required?
2. **WHAT** was authorized?
3. **WHO** possessed the authority?
4. **WHAT** boundary applied?
5. **WAS** the authority valid when relevant action occurred?
6. **WAS** the boundary technically and procedurally enforceable?
7. **WHAT** was actually implemented?
8. **WHAT** actually operated?
9. **WHAT** evidence demonstrates that operation?
10. **WAS** the evidence sufficiently reliable and complete for the assertion?
11. **COULD** unauthorized or prohibited paths materially bypass the control?
12. **COULD** the action be revoked where revocation was required?
13. **DID** material changes invalidate prior assumptions or assurance?
14. **IS** accountability clearly established?
15. **IS** the conclusion within the defined assurance boundary?
16. **CAN** the conclusion withstand credible contradictory evidence?
17. **HAVE** applicable legal, regulatory, contractual, and professional requirements been separately considered?

49. GOVERNANCE DRIFT CONE

The Framework can be represented as:

GOVERNANCE INTENT



APPLICABLE CRITERIA



AUTHORITY



BOUNDARY



ENFORCEMENT



IMPLEMENTATION



OPERATION



EVIDENCE



TESTING




REVOCATION



ASSURANCE



CONTINUOUS REVALIDATION



At every transition, the assurance question is:

HAS THE GOVERNANCE STATE REMAINED CONSISTENT WITH THE APPLICABLE CRITERIA, AUTHORITY, BOUNDARY, CONTROL OBJECTIVE, EVIDENCE REQUIREMENTS, AND ASSURANCE SCOPE?

50. GOVERNING PRINCIPLE

GOVERNANCE IS NOT PROVEN BY DOCUMENTATION.

Documentation establishes intent or prescribed requirements.

Authorization establishes permitted authority.

Implementation establishes what was built or configured.

Operation establishes what occurred.

Evidence establishes what can be demonstrated.

Testing establishes what was evaluated.

Continuing revalidation establishes whether prior assurance remains reliable when relevant conditions change.

Therefore:

NO GOVERNANCE CLAIM SHALL EXCEED THE APPLICABLE CRITERIA, AUTHORITY, EVIDENCE, TESTING, SCOPE, PERIOD, LIMITATIONS, OR PROFESSIONAL BASIS SUPPORTING THAT CLAIM.

51. FINAL POSITION

The Framework's central proposition is:

AN AGENTIC AI SYSTEM'S TECHNICAL ABILITY TO PERFORM AN ACTION DOES NOT ESTABLISH AUTHORITY TO PERFORM THAT ACTION. VALID AUTHORIZATION DOES NOT, BY ITSELF, ESTABLISH CORRECT IMPLEMENTATION, EFFECTIVE ENFORCEMENT, ACTUAL OPERATION, CONTINUING CONFORMITY, OR EFFECTIVE REVOCATION. THOSE PROPOSITIONS REQUIRE EVIDENCE AND TESTING APPROPRIATE TO THE APPLICABLE ASSERTION, CRITERIA, RISK, SCOPE, PERIOD, AND ASSURANCE OBJECTIVE.

Where material conditions can change:

PRIOR ASSURANCE SHALL NOT BE ASSUMED TO REMAIN VALID WITHOUT CONSIDERATION OF WHETHER THE CONDITIONS SUPPORTING THAT ASSURANCE HAVE REMAINED SUBSTANTIALLY UNCHANGED.

52. FINAL REGULATORY, LEGAL & PROFESSIONAL SAFEGUARD

THIS FRAMEWORK IS A METHODOLOGY — NOT A LAW, REGULATION, LEGAL OPINION, REGULATORY DETERMINATION, PROFESSIONAL STANDARD, CERTIFICATION, OR SUBSTITUTE FOR ANY APPLICABLE AUTHORITATIVE REQUIREMENT.

Its purpose is to provide a disciplined and jurisdiction-neutral method for examining Agentic AI governance, authority, boundaries, implementation, enforcement, operation, evidence, testing, revocation, accountability, and continuing assurance.

Where an applicable authoritative requirement imposes a higher, different, or more specific obligation:

THAT REQUIREMENT GOVERNS.

Where uncertainty exists regarding legal, regulatory, contractual, or professional applicability, the matter shall be appropriately escalated rather than resolved by assumption.

53. FINAL ASSURANCE STATEMENT

GOVERNANCE IS DEMONSTRATED ONLY TO THE EXTENT THAT THE RELEVANT ASSERTION IS SUPPORTED BY APPROPRIATE CRITERIA, SUFFICIENT AND APPROPRIATE EVIDENCE, APPROPRIATE TESTING, DEFINED SCOPE, DEFINED PERIOD, IDENTIFIED LIMITATIONS, AND ACCOUNTABLE AUTHORITY.

NO FRAMEWORK, CONTROL, TEST, OR ASSURANCE CONCLUSION SHOULD BE REPRESENTED AS ELIMINATING LEGAL, REGULATORY, PROFESSIONAL, OPERATIONAL, OR OTHER RESIDUAL RISK.

WHERE MATERIAL CONDITIONS CAN CHANGE, THE ORGANIZATION SHOULD BE CAPABLE OF IDENTIFYING WHEN PRIOR ASSURANCE MAY NO LONGER BE RELIABLE AND OF REVALIDATING THAT ASSURANCE AS APPROPRIATE.

NO CLAIM SHALL EXCEED THE AUTHORITY, CRITERIA, EVIDENCE, TESTING, SCOPE, PERIOD, LIMITATIONS, OR PROFESSIONAL BASIS SUPPORTING IT.

CORE FORMULA

INTENT

- CRITERIA
- AUTHORITY
- BOUNDARY
- ENFORCEMENT
- IMPLEMENTATION
- OPERATION
- EVIDENCE
- TESTING

- REVOCATION
- ASSURANCE
- REVALIDATION

ULTIMATE PRINCIPLE

PROVE THE GOVERNANCE STATE — DO NOT ASSUME IT.

PROVE THE AUTHORITY — DO NOT INFER IT.

PROVE THE CONTROL — DO NOT EQUATE DOCUMENTATION WITH EFFECTIVENESS.

PROVE THE EVIDENCE — DO NOT CONFUSE ABSENCE OF EVIDENCE WITH EVIDENCE OF ABSENCE.

PROVE THE ASSURANCE BOUNDARY — DO NOT OVERSTATE THE CONCLUSION.

AND WHERE AN AUTHORITATIVE LEGAL, REGULATORY, CONTRACTUAL, OR PROFESSIONAL REQUIREMENT APPLIES, LET THAT REQUIREMENT GOVERN.

